

ThinkTwice: Towards Scalable Decoders for End-to-End Autonomous Driving

Authors: Xaosong Jia, Penghao Wu, Li Chen, et al.

Institutions: Shanghai Jiao Tong University, Shanghai AI Laboratory, University of California at San Diego



M11352023 徐彥歲

2024/11/19

Agenda

- Introduction and Background
- Research Objectives
- ThinkTwice Methodology
 - Overall Architecture
 - Key Modules: Look, Prediction, Refinement
- Experiments and Results
- Contributions and Future Directions
- Conclusion

Introduction & Background

End-to-End Autonomous Driving

- Maps raw sensor data (e.g., cameras, LiDAR) to control signals or future trajectories.
- Efficient and avoids error accumulation in traditional modular pipelines.


Challenges in Current Methods

- Encoder-Decoder Paradigm:
 - Encoder: Extracts features from sensor data.
 - Decoder: Predicts future trajectories or actions.
 - Imbalance: Encoders are sophisticated (e.g., ResNet, Transformer), while decoders are simple (e.g., MLPs, GRUs).
- Burden on Decoders:
 - Identify safety-critical regions from massive input fields.
 - Infer future situations accurately.

Research Objectives

- 1. Fully utilize the encoder's capacity by conditioning it on future predictions.**
- 2. Increase decoder capacity to handle spatial-temporal reasoning effectively.**


ThinkTwice Architecture



Overview:

- Input: Multi-modal sensor data(cameras, LiDAR).
- Output: Accurate trajectories and control actions.
- Modules:
 - Coarse Prediction
 - Look Module
 - Prediction Module
 - Refinement Module

ThinkTwice Architecture



Overview:

- **Input:** High-dimensional features extracted from multi-modal sensor data (e.g., cameras, LiDAR).
- **Output:** Structured intermediary representations for coarse trajectory prediction.
- **Key Processes:**
 - **Feature Conditioning:** Applies spatial-temporal priors to encoder outputs.
 - **Hierarchical Data Flow:** Organizes raw sensory features into actionable predictions.
 - **Intermediate Validation:** Ensures refined information before reaching the coarse prediction module.

Coarse Prediction Module



Role: Initial trajectory and action prediction based on encoder features.

Process:

- Encoder extracts Bird's Eye View (BEV) features.
- Multi-Layer Perceptron (MLP) generates coarse trajectory and action.
- Output serves as input to subsequent modules.

Core Mathematical Equations in the Coarse Prediction Module

1. BEV Feature Map from the Encoder:

$$\mathbf{H}_{\text{BEV}} \in \mathbb{R}^{C \times BH \times BW},$$

- C : Number of channels in the feature map.
- BH, BW : Height and width of the BEV grid.

2. Environment Feature Processing and Flattening:

$$\mathbf{H}_{\text{env}} = \text{Flatten}(\text{Conv2D}(\mathbf{H}_{\text{BEV}})),$$

- Applies a 2D convolution to downsample \mathbf{H}_{BEV} , then flattens the result into a 1D vector \mathbf{H}_{env} .
- \mathbf{H}_{env} : Encodes the environmental information from cameras and LiDAR.

3. Routing Information Encoding:

$$\mathbf{H}_{\text{mst}} = \text{MLP}(\text{Routing Information}),$$

- Encodes routing information such as the target point, high-level commands (e.g., go straight, turn left), and current speed into a compact vector.

4. Initial Action and Trajectory Prediction:

$$\mathbf{Ctrl}_0, \mathbf{Traj}_0 = \text{MLP}([\mathbf{H}_{\text{env}}; \mathbf{H}_{\text{mst}}]),$$

- Combines \mathbf{H}_{env} and \mathbf{H}_{mst} to predict the initial action \mathbf{Ctrl}_0 and trajectory \mathbf{Traj}_0 .
- \mathbf{Ctrl}_0 : Initial action prediction.
- \mathbf{Traj}_0 : Initial trajectory prediction.

Look Module




Figure 3. Overall architecture of Look Module.

Purpose: Mimics human drivers' behavior by checking target regions for safety.

Steps:

1. Identify critical regions based on predicted trajectories.
2. Retrieve fine-grained features from cameras and LiDAR.
3. Integrate features into enhanced representation.

Core Mathematical Equations in the Look Module

1. Image Feature Retrieval with Deformable Attention:

$$\mathbf{H}_{i+1}^{\text{img-look}} = \text{DeformAttn}(\mathbf{H}_{\text{Img}}; \mathbf{Traj}_i; \mathbf{H}_{\text{env}}, \mathbf{H}_{\text{mst}}),$$

- \mathbf{H}_{Img} : The multi-scale image feature maps.
- \mathbf{Traj}_i : The trajectory from the previous layer, serving as the reference point for deformable attention.
- \mathbf{H}_{env} : The environment feature, used as the query in the attention mechanism.
- \mathbf{H}_{mst} : The multi-scale transformer feature, also used as the query.

2. LiDAR Feature Retrieval:

$$\mathbf{H}_{i+1}^{\text{lidar-look}} = \text{MLP}(\text{Flatten}(\text{Voxels around } \mathbf{Traj}_i)),$$

- Retrieves voxel features surrounding the predicted trajectory \mathbf{Traj}_i .
- Applies a flattening operation followed by an MLP to produce the LiDAR-based Look feature $\mathbf{H}_{i+1}^{\text{lidar-look}}$.

3. Fusion of Image and LiDAR Look Features:

$$\mathbf{H}_{i+1}^{\text{look}} = \text{MLP}([\mathbf{H}_{i+1}^{\text{img-look}}; \mathbf{H}_{i+1}^{\text{lidar-look}}]),$$

- Concatenates the image-based Look feature $\mathbf{H}_{i+1}^{\text{img-look}}$ and the LiDAR-based Look feature $\mathbf{H}_{i+1}^{\text{lidar-look}}$.
- Passes the concatenated features through an MLP to produce the final Look feature $\mathbf{H}_{i+1}^{\text{look}}$.

4. Environment Feature Update:

$$\mathbf{H}_{\text{env}} = \text{MLP}(\mathbf{H}_{i+1}^{\text{look}}),$$

- Updates the environment feature \mathbf{H}_{env} using the final Look feature $\mathbf{H}_{i+1}^{\text{look}}$.

Prediction Module




Figure 4. Overall architecture of Prediction Module.

Purpose: Anticipates future scenarios conditioned on the vehicle's predicted actions.

Process:

1. Uses spatial GRU to model the future environment.
2. Predicts agent interactions and potential collisions.
3. Supervises future scene features via teacher forcing.

Core Mathematical Equations in the Prediction Module

Output of the Spatial GRU:

$$\mathbf{H}_{i+1}^{\text{predict}} = \text{SpatialGRU}(\mathbf{H}_{\text{BEV}}; \mathbf{Ctrl}_i^t, \mathbf{Traj}_i^t),$$

- \mathbf{H}_{BEV} : The current BEV feature, used as the initial state.
- \mathbf{Ctrl}_i^t and \mathbf{Traj}_i^t : The predicted action and trajectory from the previous layer, used as input at each time step.

Teacher Forcing:

$$\mathbf{H}_{i+1}^{\text{predict,gt}} = \text{SpatialGRU}(\mathbf{H}_{\text{BEV}}; \mathbf{Ctrl}^{\text{gt}}, \mathbf{Traj}^{\text{gt}}),$$

- Uses the ground truth action $\mathbf{Ctrl}^{\text{gt}}$ and trajectory $\mathbf{Traj}^{\text{gt}}$ as input to generate $\mathbf{H}_{i+1}^{\text{predict,gt}}$.

Shape of the Output:

$$\mathbf{H}_{i+1}^{\text{predict}} \in \mathbb{R}^{T \times BH \times BW \times C},$$

- T : The prediction horizon (number of future time steps).
- BH, BW : The height and width of the BEV grid.
- C : The hidden feature dimension.

Supervision Signal:

$$\mathcal{L}(\mathbf{H}_{i+1}^{\text{predict}}, \mathbf{H}_{i+1}^{\text{predict,gt}}),$$

- The loss function compares $\mathbf{H}_{i+1}^{\text{predict}}$ to the ground truth future scene feature $\mathbf{H}_{i+1}^{\text{predict,gt}}$, using Roach BEV features as the target for supervision.

Refinement Module

Role: Refines coarse predictions using insights from Look and Prediction Modules.

Process:

1. Calculates offset between coarse predictions and ground truth.
2. Updates actions and trajectories iteratively.
3. Enables stacking for enhanced accuracy.

Core Mathematical Equations in the Refinement Module

1. Offset Calculation:

$$\mathcal{O}_{i+1}^{\text{ctrl}}, \mathcal{O}_{i+1}^{\text{traj}} = \text{MLP}([\mathbf{H}_{i+1}^{\text{look}}; \mathbf{H}_{i+1}^{\text{predict}}; \mathbf{Ctrl}_i; \mathbf{Traj}_i; \mathbf{H}_{\text{env}}; \mathbf{H}_{\text{mst}}]),$$

- $\mathcal{O}_{i+1}^{\text{ctrl}}$: Offset for the action prediction.
- $\mathcal{O}_{i+1}^{\text{traj}}$: Offset for the trajectory prediction.
- The offsets are computed using an MLP, based on the Look feature $\mathbf{H}_{i+1}^{\text{look}}$, Prediction feature $\mathbf{H}_{i+1}^{\text{predict}}$, current action \mathbf{Ctrl}_i , current trajectory \mathbf{Traj}_i , environment feature \mathbf{H}_{env} , and routing feature \mathbf{H}_{mst} .

2. Supervision Signals:

$$\mathcal{O}_{i+1}^{\text{ctrl}} = \mathcal{L}(\mathbf{Ctrl}_i, \mathbf{Ctrl}^{\text{gt}}), \quad \mathcal{O}_{i+1}^{\text{traj}} = \mathcal{L}(\mathbf{Traj}_i, \mathbf{Traj}^{\text{gt}}),$$

- The loss function \mathcal{L} is the Smooth L1 Loss, which supervises the predicted offsets for both action and trajectory by comparing them to the ground truth.

3. Action and Trajectory Update:

$$\mathbf{Ctrl}_{i+1} = \mathbf{Ctrl}_i + \mathcal{O}_{i+1}^{\text{ctrl}}, \quad \mathbf{Traj}_{i+1} = \mathbf{Traj}_i + \mathcal{O}_{i+1}^{\text{traj}}.$$

- The predicted action and trajectory are updated by adding the offsets $\mathcal{O}_{i+1}^{\text{ctrl}}$ and $\mathcal{O}_{i+1}^{\text{traj}}$, respectively.

Experiments and Results

Setup:

- Platform: CARLA Simulator.
- Benchmarks: Town05 Long, Longest6.
- Metrics:
 - Route Completion (RC): Percentage of completed route.
 - Infraction Score (IS): Measures safety violations.
 - Driving Score (DS): Combines RC and IS.

Method	Encoder	Decoder	Modality	Extra Supervision	DS↑	RC↑	IS↑
CILRS [16]	ResNet + Flatten	MLP	C1	None	7.8±0.3	10.3±0.0	0.75±0.05
LBC [10]	ResNet + Flatten	MLP	C3	Expert	12.3±2.0	31.9±2.2	0.66±0.02
Transfuser [13]	Fusion via Transformer	GRU	C3L1	Dep+Seg+Map+Box	31.0±3.6	47.5±5.3	0.77±0.04
Roach [87]	ResNet + Flatten	MLP	C1	Expert	41.6±1.8	96.4±2.1	0.43±0.03
LAV [9]	PointPainting	Multi-layer GRUs	C4L1	Expert+Seg+Map+Box	46.5±2.3	69.8±2.3	0.73±0.02
TCP [78]	ResNet + Flatten	GRU	C1	Expert	57.2±1.5	80.4±1.5	0.73±0.02
ThinkTwice	Geometric Fusion in BEV	Look-Predict-Refine	C4L1	Expert+Dep+Seg+Map	65.0±1.7	95.5±2.0	0.69±0.05
MILE*† [29]	ResNet + Flatten	GRU	C1	Map+Box	61.1±3.2	97.4±0.8	0.63±0.03
Interfuser* [66]	Fusion via Transformer	Transformer + GRU	C3L1	Map+Box	68.3±1.9	95.0±2.9	-
ThinkTwice*	Geometric Fusion in BEV	Look-Predict-Refine	C4L1	Expert+Dep+Seg+Map	70.9±3.4	95.5±2.6	0.75±0.05

↑ means the higher the better. * denotes using extra data. † denotes no scenarios are involved in the evaluation, which is a much easier benchmark.

Experiments and Results

Setup:

- Platform: CARLA Simulator.
- Benchmarks: Town05 Long, Longest6.
- Metrics:
 - Route Completion (RC): Percentage of completed route.
 - Infraction Score (IS): Measures safety violations.
 - Driving Score (DS): Combines RC and IS.

Method	DS↑	RC↑	IS↑
WOR [8]	23.6	52.3	0.59
LAV [9]	34.2	73.5	0.53
Transfuser [13, 57]	56.7	92.3	0.62
ThinkTwice	61.3	73.0	0.81
ThinkTwice*	66.7	77.2	0.84

Table 2. Performance on Longest6 benchmark.

ID	Method	DS↑	RC↑	IS↑
1	Baseline	43.0±3.9	74.2±3.2	0.56±0.05
2	+Segmentation & Depth Task	45.0±2.2	73.2±4.0	0.60±0.04
3	+Segmentation & Depth Projection	57.8±1.5	78.8±3.3	0.74±0.03
4	+Two Frames	59.3±1.4	80.2±3.6	0.74±0.04
5	Expert Feature -> BEV Segmentation	50.5±2.3	73.4±3.2	0.70±0.03

Table 3. Ablation for Encoder Design. Baseline means original

ID	Method	DS↑	RC↑	IS↑
4	TCP-Head	59.3±1.4	80.2±3.6	0.74±0.04
5	+1 Decoder	61.6±1.4	90.8±2.2	0.68±0.04
6	+5 Decoders	65.0±1.7	95.5±2.0	0.69±0.05
7	w/o Look	62.4±1.8	93.5±2.1	0.67±0.08
8	w/o Predict	61.7±1.9	96.2±3.0	0.63±0.03
9	w/o TF	62.0±2.2	97.1±2.9	0.62±0.04

Table 4. Ablation for Decoder Design. The baseline is Model4

Talk About

01

Contributions

02

Future Directions

03

Conclusion

Thank you For Attention